

ATOC7500 – Application Lab #5

Filtering Timeseries

in class Wednesday November 11 and Monday November 16

Notebook #1 – ATOC7500_applicationlab5

[ATOC7500_applicationlab5_check_python_convolution.ipynb](#)

LEARNING GOAL

1) Understand what is happening “under the hood” in different python functions that are used to smooth data in the time domain.

Use this notebook to understand the different python functions that can be used to smooth data in the time domain. Compare with a “by hand” convolution function. Look at your data by printing its shape and also values. Understand what the python function is doing, especially how it is treating edge effects.

Notebook #2 – Filtering Synthetic Data

[ATOC7500_applicationlab5_synthetic_data_with_filters.ipynb](#)

LEARNING GOALS:


- 1) Apply both non-recursive and recursive filters to a synthetic dataset
- 2) Contrast the influence of applying different non-recursive filters including the 1-2-1 filter, 1-1-1 filter, the 1-1-1-1-1 filter, and the Lanczos filter.
- 3) Investigate the influence of changing the window and cutoff on Lanczos smoothing.

DATA and UNDERLYING SCIENCE:

In this notebook, you analyze a timeseries with known properties. You will apply filters of different types and assess their influence on the resulting filtered dataset.

Questions to guide your analysis of Notebook #2:





- 1) Create a red noise timeseries with oscillations. Plot your synthetic data – Look at your data!! Look at the underlying equation. What type of frequencies might you expect to be able to remove with filtering?




We introduce 2 different frequencies to our red noise time that we could potentially remove: $52/256$ (low frequency) and $100/256$ (high frequency).

2) Apply non-recursive filters in the time domain (i.e., apply a moving average to the original data) to reduce power at high frequencies. Compare the filtered time series with the original data (top plot). Look at the moving window weights (bottom plot). You are using the function “`filtfilt`” from `scipy.signal`, which applies both a forward and a backward running average. Try different filter types – What is the influence of the length of the smoothing window or weighted average that is applied (e.g., 1-1-1 filter vs. 1-1-1-1-1 filter)? What is the influence of the amplitude of the smoothing window or the weighted average that is applied (e.g., 1-1-1 filter vs. 1-2-1 filter)? Tinker with different filters and see what the impact is on the filtering that you obtain.

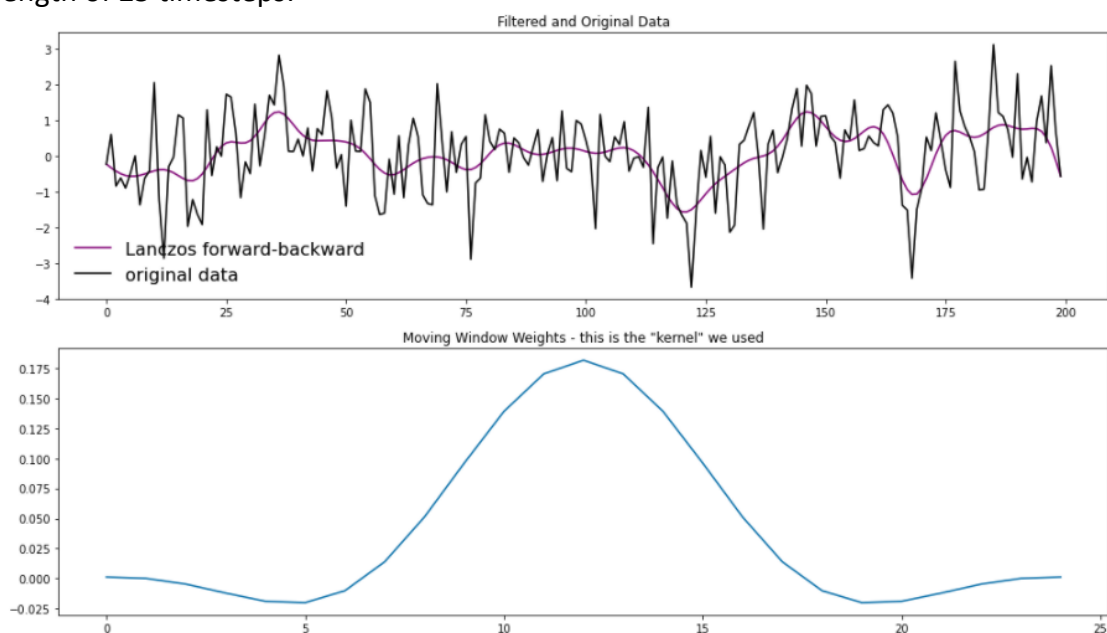
A 1-1-1 filter creates a more smoothed time series than a 1-2-1 filter because it puts more weight on the surrounding points and less weight on the individual points. A 1-1-1-1-1 filter creates even smoother data because it considers a larger number of surrounding values.



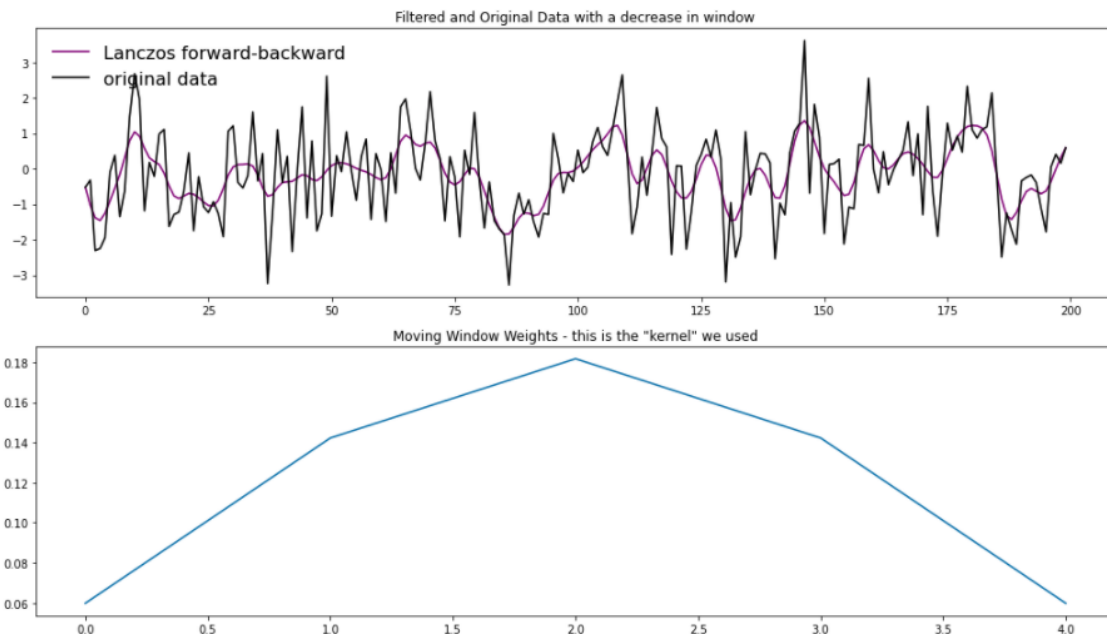


3) Apply a Lanczos filter to remove high frequency noise (i.e., to smooth the data). What is the influence of increasing/decreasing the window length on the smoothing and the response function (Moving Window Weights) in the Lanczos filter? What is the influence of increasing/decreasing the cutoff on the smoothing and the response function?


Below is the result of using a Lanczos filter to remove high frequency noise with a window length of 25 timesteps:



If we decrease the window length to 5 timesteps, we can see that the resulting time series is not as smooth.




If we increase the cutoff, more weight is concentrated in the middle of the window and the resulting tie series is not smoothed as much. We can see this effect in the plot below:



4) Apply a Butterworth filter, a recursive filter. Compare the response function (Moving Window Weights) with the non-recursive filters analyzed above.

With a Butterworth filter we can specify an exact frequency we want to cutoff. In comparison, with the other non-recursive filters we analyzed above we can experiment with different kernels to try and eliminate certain frequencies but we can't specify the exact frequencies we wish to eliminate. For example, below we see the result of 2 different butterworth filters. The first eliminates all frequencies greater than 0.25 and the second all greater than 0.5. We can see that the first filter creates a smoother time series.



Notebook #3 – Filtering ENSO data

[ATOC7500_applicationlab5_mrbutterworth_example.ipynb](#)

LEARNING GOALS:

- 1) Assess the influence of filtering on data in both the time domain (i.e., in time series plots) and the spectral domain (i.e., in plots of the power spectra).
- 2) Apply a Butterworth filter to remove power of specific frequencies from a time series.
- 3) Contrast the influence of differing window weights on the filtered dataset both in the time domain and the spectral domain.
- 4) Calculate the response function using the Convolution Theorem.
- 5) Assess why the python function `filtfilt` is filtering twice.


DATA and UNDERLYING SCIENCE:

In this notebook, you analyze monthly sea surface temperature anomalies in the Nino3.4 region from the Community Earth System (CESM) Large Ensemble project fully coupled 1850 control run (<http://www.cesm.ucar.edu/projects/community-projects/LENS/>). A reminder that an pre-industrial control run has perpetual 1850 conditions (i.e., they have constant 1850 climate). The file containing the data is in netcdf4 format: CESM1_LENS_Coupled_Control.cvdv_data.401-2200.nc

Does this all look and sound really familiar? It should!! This dataset is the same one you analyzed in Homework #4.


Questions to guide your analysis of Notebook #3:

- 1) Look at your data! Read in your data and Make a plot of your data. Make sure your data are anomalies (i.e., the mean has been removed). Look at your data. Do you see variance at frequencies that you might be able to remove?




There seems to be some high frequency variance (a couple months to a year) and some low frequency variance (tens of years) that I might be able to remove

- 2) Calculate the power spectrum of your original data. Calculate the power spectra of the Nino3.4 SST index (variable called “nino34”) in the fully coupled model 1850 control run. Apply the analysis to the first 700 years of the run. Use Welch’s method (WOSA!) with a Hanning window and a window length of 50 years. Make a plot of normalized spectral power vs. frequency. Where is their power that you might be able to remove with filtering?




There is power from approximately 0.01 to 0.07. We might be able to remove power from some of these frequencies by filtering.

3) Apply a Butterworth Filter. Use a Butterworth filter to remove all spectral power at frequencies greater than 0.04 per month (i.e., less than 2 year). Use an order 1 Butterworth filter ($N=1$, 1 weight). Replot the original data and the filtered data. Calculate the power spectra of your filtered data. Assess the influence of your filtering in both in time domain (i.e., by comparing the original data time series and filtered time series data) and the frequency domain (i.e., by comparing the power spectrum of the original data and the power spectrum of the filtered data). Look at the response function of the filter in spectral domain using the convolution theorem. Well that was pretty boring... we still have most of the power retained....




4) Let's apply another Butterworth Filter and this time really get rid of ENSO power!. Let's really have some fun with the Butterworth filter and have a big impact on our data... Let's remove ENSO variability from our original timeseries. Apply the Butterworth filter but this time change the frequency that you are cutting off to 0.01 per month (i.e., remove all power with timescales less than 8 years). Use an order 1 filter ($N=1$). Replot the original data and the filtered data. Calculate the power spectra of your filtered data. Assess the influence of your filtering in both in time domain (i.e., by comparing the original data time series and filtered time series data) and the frequency domain (i.e., by comparing the power spectrum of the original data and the power spectrum of the filtered data). Look at the response function of the filter in spectral domain using the convolution theorem.



5) Let's apply yet another Butterworth Filter – and this time one with more weights. Repeat step 4) but this time change the order of the filter. In other words, increase the number of weights being used in the filter by increasing the parameter N in the jupyter notebook. What is the impact of increasing N on the filtered dataset, the power spectra, and the moving window weights? You should see that as you increase N – a sharper cutoff in frequency space occurs in the power spectra. Why?

As N increases we put more power on the central frequency of the window which leads to a sharper, less smooth cutoff in the frequency space of the power spectra.



6) Assess what is “under the hood” of the python function. How are the edge effects treated? Why is the function `filtfilt` filtering twice?

There are different options for handling the edge effects: 1) `pad`, which pads the ends of the signal or 2) `'gust'`, which uses Gustaffson’s method. The `filtfilt` function filters twice so that the resulting signal has zero phase.